

Matlab Tutorial for COMP24111

(includes exercise 1)

1 Exercises to be completed by end of lab

There are a total of 11 exercises through this tutorial. By the end of the lab, you should have completed the first 10 exercises each worth 1%. If you have completed everything else, you are very welcome to keep challenging Exercise 11, which is optional and will not be marked. Your attendance will not be signed off unless you are seen by a TA to be making an effort.

2 Submission instructions

Your work should be in a directory `COMP24111/ex1/`. You should submit in the usual manner, using the `submit` command from within that directory. This week you need to submit just 1 file :

`ex1code.zip` — a zip file containing all of your matlab code

It is recommended to include your matlab code for the 10 exercises in one single “.m” file. In this case, your .zip file will only contain one .m file.

The deadline for submitting your code is the end of your first lab. Also, you will need to get the exercise marked (out of 10) by a TA during your first lab.

3 What is Matlab

Matlab is a software package which was developed for numerical analysis involving matrices (“matlab” as in **matrix laboratory**). It is now used for general technical computing. It is an interpreted language, which means that commands are run as they are entered in, without being compiled first. It also means that most of the commands, except the most basic ones, are defined in text files which are written in this interpreted language. So, you can read these files, and add to the language by making similar files. Many commands you will use have been created specifically for these labs.

Matlab is an imperative language and is like C in several respects. The basic data element for Matlab is the matrix or array. This means that you can write programs which act on arrays easily and without the need for dimensioning and memory allocation. Because it is interpreted, it is slower than C, and to get the fastest performance, the matrix nature of Matlab must be used fully (the programs must be “vectorized”).

Matlab has built-in graphics and visualization tools. There are many add-on “toolboxes”. Matlab is an excellent prototyping language, because it has so many useful mathematics utilities, built-in. Many of the latest algorithms and research ideas in machine learning appear as Matlab packages before they are produced in other forms, such as C++.

When one get experience with the software, one can produce algorithms in matlab much more quickly than one could in JAVA, say. The downside is that these will run much more slowly than if they were written in C++, or even in JAVA. This is why Matlab is often used to prototype ideas. When the algorithms applied to large systems and need to run fast, they are often rewritten in a compiled language, such as C or C++.

4 Starting Matlab

Matlab **runs on Linux machines in the computer science department**. It also runs under Windows. For the lab exercises, you should run it under Linux. You start it just by typing at the Linux command line:

```
matlab
```

and pressing enter.

5 Things to know about Matlab

Matlab can process Unix commands. Type `ls` at the Matlab command line to see your files. Matlab does makes some changes, for example `rm` is replaced by `delete`. However you can ‘drop-down’ to the full unix terminal by including an exclamation mark `!` at the start of the line, for example `!rm` will the unix command.

Matlab has a built-in editor. Type `edit` at the command line to start it. You can then save your code in a file with extension `.m` and run it. If your file is called `ex1.m` you run it just by typing `ex1` at the Matlab prompt.

6 The Matlab Environment

6.1 The Matlab Environment

The Matlab environment consists of a window with three sub windows (fig 1).

The Command Window: This is the most important window. This is where the commands to Matlab will be entered. All of the commands described below are run by typing them at the prompt `>>` in the Command Window.

The Workspace Window: This shows the variables that currently exist in the workspace. Matlab is an interpreted language, and variables which you define and use remain until they are deleted or until Matlab is exited. You can click on the square icons to see their values.

Command History Window: This shows all of the commands you have ever entered. To rerun commands, you can cut and paste them from the Command History window to the Command window.

There is also a tab to a window which shows the files in the current working directory (the directory from which Matlab was invoked). You can remove all but the Command Window from the View menu on the menu bar.

Start Matlab. If you want to see some of the things which can be done in Matlab, type `demos` at the prompt, and run some demos. Other Demos under menus Matlab/Demos is pretty good. Otherwise, carry on with the tutorial.

7 Getting Help

If you need more information about any Matlab function, there are several ways of getting it:

1. At the prompt, type `help` followed by the function name, e.g.

```
>> help sum
```



(type `'help'` on its own to get a list of help topics.) N.B. Matlab online help entries use uppercase characters for the function and variable names to make them stand out from the rest of the text. When typing function names, however, always use the corresponding lowercase characters because Matlab is case sensitive and most function names are actually in lowercase.

2. The Help menu from the menu bar gives access to a huge range of documents and tutorials. Look under “MATLAB Help”. “Getting Started” which contains a good tutorial. “Using Matlab” is a useful reference.

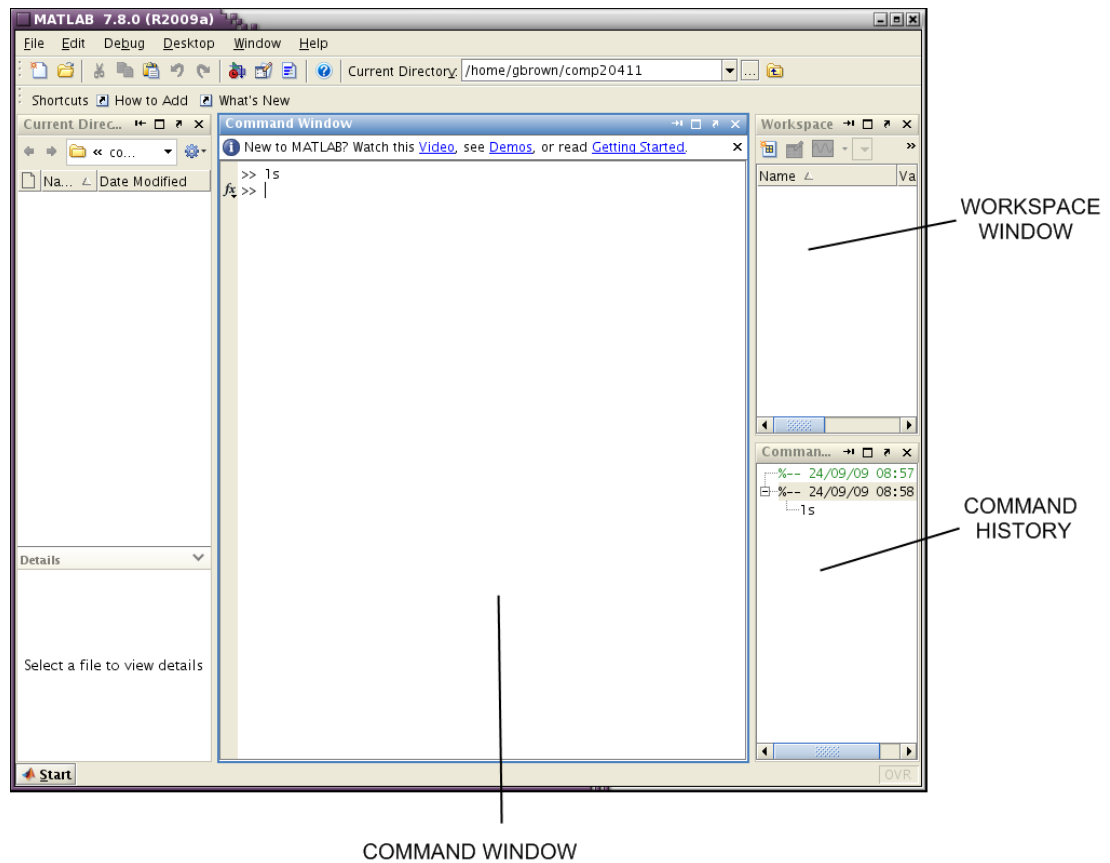


Figure 1: The Matlab Window. The Command Window is where you enter commands. The Workspace Window shows the variables defined in the workspace. The Command History window stores a history of the commands you have run. You can move this child windows around within the main Matlab window if you wish.

8 Entering Commands

In Matlab, commands are entered at the prompt, and run when the return key is entered. For example, if you wanted to know the value of 2π , you could enter

```
x=2*pi
```

at the prompt (**pi is a built-in constant**). Matlab will give the answer,

```
x =
```

```
6.2832
```

and create a variable x in the workspace (if it did not already exist), and set its value to the above. If you put a semicolon after the command,

```
x=2*pi;
```

the answer will not be printed on the screen, but the value of x will be set. This is useful when you don't need to see the value (because it is part of an intermediate calculation), or when the variable being set is a large structure which would take many pages to print out. If you don't give a variable name to the calculation, the variable is stored in a variable called **ans**. E.g.,

```
2*pi
```

results in

```
ans =
```

```
6.2832
```

Exercise 1

Use Matlab to calculate 15 factorial (i.e. $15!$). Use help to find the command and how to use it, then type it in to get the answer.

9 Matrices

One of the most important aspects of Matlab how it handles matrices. Whereas other programming languages work with numbers one at a time, Matlab allows you to work with entire matrices, which can take some getting used to. A matrix is a rectangular array of numbers, like a two-dimensional array in C or JAVA. An “ n by m ” matrix has n rows and m columns. Special meaning is sometimes attached to 1 by 1 matrices, which are called “scalars” (ordinary numbers, basically), and to matrices with only one row or only one column, which are called “vectors”.

9.1 Entering Matrices

There are several ways to enter matrices in Matlab. These include:

- Entering an explicit list of elements.
- Loading matrices from external data files.
- Generating matrices using functions.

To enter a matrix explicitly, there are a few basic rules to follow:

- Separate the elements of a row with blanks or commas.
- Use a semicolon, ; or carriage returns, to indicate the end of each row.
- Surround the entire list of elements with square brackets, [].

For example, to input a 4×4 magic square enter:

```
>> A = [16 3 2 13; 5 10 11 8; 9 6 7 12; 4 15 14 1]
```

at the prompt. This describes the matrix

16	3	2	13
5	10	11	8
9	6	7	12
4	15	14	1

and assigns it to the variable 'A'. If you don't know what a magic square is, this will be explained later in this document (exercise 6). **Be careful - Matlab is case-sensitive, so it distinguishes between 'A' and 'a'.**)

Matlab will echo the matrix back at you, unless you put a semi-colon (;) at the end of the line. This is very useful when the matrices are very large.

This matrix can be referred to as A until the end of the session, unless you decide to alter it in some way. **When it encounters a variable it hasn't seen before, Matlab automatically creates one.** To see which variables have been created so far, look in the Workspace submenu. If you type A (or any defined variable name) at the Matlab prompt, it will print back the name with its contents. This is useful if you want to check the contents of a variable.

There are commands for creating special matrices. These take parameters which define the size of matrix they generate. Some of these are:

zeros: makes a matrix of all zeros (for initialization, for example). For example, `zeros(2,3)` makes a 2 by 3 matrix of zeros.

ones: makes a matrix of all ones. Used like the zeros command.

eye: makes an identity matrix. E.g. `eye(10)` makes 10 by 10 matrix with 1's on the diagonal and 0's off the diagonal.

Exercise 2

Enter the **magic square A** (further up this page). Set B to be a 4 by 4 matrix with 1's on the diagonal and 0's elsewhere. Check to see that both variables are present.

9.2 Accessing Elements of a Matrix

An element of a matrix can be referred to by using the format $M(\text{row}, \text{column})$. For example, $A(3,2)$ is 6. So to calculate the sum of the top row of A , one could use

```
A(1,1)+A(1,2)+A(1,3)+A(1,4)
```

(there simpler ways to do this, as shown in the next section).

A range of the array can be referred to by using the colon operator. $M(i:j,k)$ refers to the rows i through j of the k th column. For example,

```
>> A(2:4,3)
```

yields,

```
ans = 11
      7
      14
```

The colon by it self refers to the entire row or column. For example, $A(:,2)$ is the entire second column of A ,

```
>> A(:,2)
```

```
ans =
```

```
3
10
6
15
```

Thus, $A(:,2)$ is equivalent to $A(1:4,2)$.

Exercise 3 Get Matlab to compute the sum of the third column of A .

9.3 More On The Colon Operator

The colon $(:)$ is one of Matlab's most important operators. It occurs in several different forms. The expression $1:10$ is a row vector containing the integers from 1 to 10 (i.e. $[1\ 2\ 3\ 4\ 5\ 6\ 7\ 8\ 9\ 10]$).

To obtain non-unit spacing, specify an increment. For example:

```
>> 100:-7:50
```

```
ans =
```

```
100    93    86    79    72    65    58    51
```

```
>> 0:pi/4:pi
```

```
ans =
    0    0.7854    1.5708    2.3562    3.1416
```

(pi is a built-in scalar constant).

Subscript expressions involving colons refer to portions of a matrix, as we have seen. So, `A(1:3,4)` means the same as `A([1,2,3],4)`,

```
>> A(1:3,4)
```

```
ans =
```

```
13
 8
12
```

```
>> A([1,2,3],4)
```

```
ans =
```

```
13
 8
12
```

Exercise 4

Generate a 4 by 2 matrix consisting of the odd-numbered columns of A using the colon operator.

9.4 Creating Random Numbers

MATLAB has built-in algorithms that make your results appear to be random and independent. The results also pass various statistical tests of randomness and independence. You can use these numbers as if they are truly random and independent. There are commands for creating random numbers. Some of these are:

rand: makes a matrix of random numbers uniformly distributed between 0 and 1. E.g. `rand(1,10)` makes a column of random numbers.

randn: as `rand`, except the numbers are normally distributed.


randi: makes a matrix of random integers from the uniform distribution in the specified range. E.g., `randi([-2,2], 1, 10)` makes a column of random integers between -2 and 2.

randperm: makes a row vector containing permuted integers. E.g., `randperm(5)` returns a row vector of a random permutation of the integers from 1 to 5. `randperm(10, 3)` returns a row vector containing 3 unique integers selected randomly from 1 to 10.

Exercise 5

Generate a 4 by 2 matrix containing two different columns that are randomly selected from the columns of **A**. Generate a 3 by 4 matrix containing three different rows that are randomly selected from the rows of **A**.

9.5 Functions on Matrices

Matlab has a number of built-in functions which can be performed on matrices. Here is a list of the more immediately useful s.

sum: Returns the sum of the columns of a matrix, or, if used on a row vector, the sum of the row. For example,

```
>> sum(A)
```

sums the columns in the matrix, and returns the result as a 1 by 4 matrix. Notice how the special variable ‘ans’ is used to store the temporary result of the operation.

mean: Returns the mean (average) of the columns of a matrix.

transpose: To return the transpose of a matrix, append an apostrophe (or ”single-quote”) to the name. For example:

```
>> A'
```

```
ans =
    16     5     9     4
     3    10     6    15
     2    11     7    14
    13     8    12     1
```

```
>> sum(A')'
```

```
ans =
    34
    34
    34
    34
```

diag: Returns the diagonal of the matrix **M**.

sqrt: Returns the square root of the elements of matrix **M**.

size: Returns the dimensions of the matrix **M**. This returns a list of two values; the form is like this

```
>> [rows columns] = size(A)
```



sort : Sorts array elements. If **M** is a vector, then **sort(M)** sorts the vector elements. If **M** is a matrix, then **sort(M)** treats the columns of **M** as vectors and sorts each column.

Exercise 6

A magic square is a matrix in which the sum of each row, each column, and each diagonal is the same. The matrix **A** is a magic square. Check that for **A**, the sum of all rows and all columns are the same, using the **sum** command and the transpose operator. Sort each row of the matrix **A** using the **sort** command and the transpose operator.

10 Operators on Matrices

You can add or subtract two matrices

```
>> A + A
```

```
ans = 32    6    4    26
       10   20   22   16
       18   12   14   24
        8   30   28    2
```

The result could easily have been stored in another variable, e.g.:

```
>> S = A + A;
```

The multiplication operator *****, division operator **/** and power operator **^** refer to *matrix* multiplication, division, and power respectively.

If a dot is put before these operators, the operator acts component by component. For example,

```
>> A.*A
```

returns the matrix containing the square of each component of **A**, whereas

```
>> A*A
```

performs matrix multiplication between the two. On scalars, both forms have the same meaning (which is the usual meaning).

```
>> A^n
```

performs matrix multiplication of **A*A*A...A**, including a total *n* matrices.

```
>> A.^n
```

performs element-wise power, by raising each element of **A** to the corresponding power *n*.

Exercise 7

Since **B** is the identity matrix, multiplication of **A** by **B** should yield **A**. Check this. What does component by component multiplication give? Check this.

Logical operations are also allowed. These are the same as in most other languages: `&`, `|`, `~`, `xor` have their usual meanings when applied to scalars. Any non-zero number represents True, and zero represents False. They can also be applied to matrices, and the result is a matrix of 0's and 1's. For example:

```
>> L = [0 0 1 1; 0 1 0 1];  
>> L(3,:) = L(1,:) & L(2,:)
```

```
L =  
    0    0    1    1  
    0    1    0    1  
    0    0    0    1
```

```
>> L(3,:) = xor(L(1,:), L(2,:))
```

```
L =  
    0    0    1    1  
    0    1    0    1  
    0    1    1    0
```

Relation operators are similar to that of other languages: `==`, `<`, `>`, `<=`, `>=`, `~=`. These return either 1 or 0, in much the same way as the logical operators:

```
>> 5<3
```

```
ans =  
    0
```

```
>> 4==2*2
```

```
ans =  
    1
```

However, these can be used with matrices as well:

```
>> A>10
```

```
ans =  
  
    1    0    0    1  
    0    0    1    0  
    0    0    0    1  
    0    1    1    0
```

Exercise 8

Use `sum` and logical operators to count the number of values of `A` which are greater than 10.

11 Graphics

11.1 Graphing functions

Matlab has range of built-in graphics and plotting routines. The command `plot(x,y)` makes a two-dimensional plot of x against y . For example,

```
x=-20:0.01:20;
y=sin(x)./x;
plot(x,y);
```

graphs the function $\sin(x)/x$ between -20 and 20 . (Try it.) The points of the x axis are separated by `0.01`; for different spacing, you would use a different increment in the colon operator in the definition of x .

Type `help plot` at the prompt to get a description of the use of the plot function. Your math teacher may have taught you to always label your axes; here is how: `xlabel` and `ylabel` puts strings on the axes, like so,

```
xlabel('x');
ylabel('sin(x)/x');
```

To get rid of the figure, type `close` at the Matlab prompt.

There are a load of options to `plot`, which `help plot` will show. It is possible to control whether the plots are lines or points, the line style, color, and other properties of the plots. The basic form is `plot(x,y,str)` where `str` is a string of one to three characters denoting color, symbol plotted at each point (if any) and line type (if any). Here is a table of options,

COLOR		POINT STYLE		LINE STYLE	
character	color	character	symbol	character	line style
b	blue	.	point	-	solid
g	green	o	circle	:	dotted
r	red	x	x-mark	-.	dashdot
c	cyan	+	plus	-	dashed
m	magenta	*	star		
y	yellow	s	square		
k	black	d	diamond		
		v	triangle (down)		
		^	triangle (up)		
		<	triangle (left)		
		>	triangle (right)		
		p	pentagram		
		h	hexagram		

For example, if we wanted to plot the graphs above as crosses, and in red, the command would be

```
plot(x,y,'r+');
```

This is useful for plotting data. For example, suppose we had some data arranged in columns.

```
data =
```

5.1000	3.5000
4.9000	3.0000
4.7000	3.2000
4.6000	3.1000
5.0000	3.6000
5.4000	3.9000
4.6000	3.4000
5.0000	3.4000
4.4000	2.9000
4.9000	3.1000

We could plot it using the command, `plot(data(:,1),data(:,2),'+')`.

Exercise 9	Plot the log of the integers from 1 to 100.
-------------------	---

11.2 Multiple Plots

If you want to compare plots of two different functions, calling `plot` twice in succession will not be satisfactory, because the second plot will overwrite the first. You can ensure a new plot window for a plot by calling `figure` first.

If you want to put multiple plots on the same figure, we set `hold on`. The default is `hold off`, which makes a new figure overwrite the current one.

Here is an example. Suppose we want to compare the log function with the square root function graphically. We can put them on the same plot. By default, both plots will appear in blue, so we will not know which is which. We could make them different colors using options. Here is the final answer,

```
x=1:100;
y=log(x);
z=sqrt(x);
plot(x,y,'r'); % plot log in red
hold on;
plot(x,z,'b'); % plot log in blue
hold off;
```

What appears after the `%` on a line is a comment. The options can also be used to plot the values as points rather than lines. For example, `'+'` plots a cross at each point, `'*'` a star and so forth. So,

```
x=1:100;
y=log(x);
```

```

z=sqrt(x);
plot(x,y,'r+'); % plot log as red crosses
hold on;
plot(x,z,'b*'); % plot log as blue stars
hold off;

```

11.3 Three-Dimensional Plots



These will not be used for this course, but you can also make three dimensional plots. If you have three dimensional data, such as

```

data =

    5.1000    3.5000    1.4000
    4.9000    3.0000    1.4000
    4.7000    3.2000    1.3000
    4.6000    3.1000    1.5000
    5.0000    3.6000    1.4000
    5.4000    3.9000    1.7000
    4.6000    3.4000    1.4000
    5.0000    3.4000    1.5000

```

This could be plotted as points in 3-d, using the `plot3` command,

```
plot3(data(:,1),data(:,2),data(:,3),'+')
```

You can rotate the plot around.

You can also plot functions of two variables. An example of this can be seen with the following

```

[x,y] = meshgrid(-8:.5:8);
r = sqrt(x.^2+y.^2) + eps;
z = sin(r)./r;
surf(x,y,z);

```

You can click on the icon which looks like a circular arrow, and rotate the plot around using the mouse. To learn about using Matlab to make all types of plots, see “Graphics” under the “Getting Starting” and under the “Using Matlab” in the on-line help. For a list of commands, type `help graph2d` or `graph3d` at the Matlab prompt.

12 Control Structures

There are two important control structures for these labs.

12.1 If

The form for **if-else** constructs is

```
if condition
    statements
else
    more statements
end
```

where the else part is optional. Likewise, there is an **elseif** keyword,

```
if condition1
    statements
elseif condition2
    more statements
elseif condition3
    even more statements
...
else
    yet more statements
end
```

12.2 For

There are loop structures using **for**. The basic construction is,

```
for index = j:k
    statements
end
```

As an example, lets sum the main diagonal of A:

```
>> c=0;
>> for i=1:4
        c=c+A(i,i);
    end;
>> c

c =

    34
>>
```

Of course, this is more easily done using `sum(diag(A))`.

Exercise 10

Use **for** loop to calculate the reciprocal of each element of A. The reciprocal of a number x is simply $1/x$. Now with a new matrix B, where each element in B is the reciprocal of the corresponding element in A, **sum the columns.** Now do this in a **single operation, using matrix commands, without a for loop.**

13 Running Commands From a File

13.1 Scripts

You can type a set of commands into a file and read that file into Matlab. Matlab will run these just as if you had typed them. Such files are called scripts. Matlab looks for files in directories defined by a `path` variable. The directory from which Matlab was invoked is in the path. To add a directory to the path, use

```
addpath directory path .
```

13.2 User-defined functions

You can create your own functions for use in Matlab. These are defined in separate files, somewhere on your search path (usually they will be in your current directory). The best way to show this is for you to try it yourself: in a text editor or in the Matlab Editor, make a file in your current directory, calling it 'test.m', with the following text in it:

```
function result = test(m)
% A test function for Matlab

result = sum(diag(m))
```

Save it, and then in Matlab type

```
>> test(A)
```

The returned value (here stored in 'ans') will be the sum of the diagonal of A. Here's a brief explanation of the file you typed in:

- 'function' tells Matlab that this is a function file
- 'result = test(m)' means that the function is called 'test', and will accept one input parameter and one output parameter, with the specified names. The file must have the same name as the function and the extension ".m".
- Anything between a % and the end of the line is treated as a comment.
- 'result = sum(diag(m))' does the actual calculation. The answer is stored in 'result', and as there are no more lines of code, the function ends. The return parameter is the last value that 'result' had.

Matlab contains a built-in editor which you can invoke with the command `edit` or by using Open or New from File on the menu bar. You can also use `textedit`, `emacs`, or whatever other UNIX editor you are familiar with.

Here is an example to test whether the sum of the diagonal from top left to bottom right of a matrix is the same as that from top right to bottom left.


```

function result = diagonalSum(m)
% DIAGONALSUM returns 1 if sum of left to right diagonal of a square
% matrix is the same as that of the right to left diagonal.

[height width]=size(m);
if (width ~=height)
    error('Only works for square matrices'); % print error message and
                                              % terminate
end
leftright=0;
rightleft=0;
for i=1:width
    leftright=leftright+m(i,i);
    rightleft=rightleft+m(i,width-i+1);
end
result=(leftright==rightleft)

```

13.3 Vectorization

Built-in matrix commands are compiled (along with the rest of Matlab) and thereby run faster. Thus, commands will be faster if they are written as built-in matrix operations. For example, the function `diagonalSum` above would run faster if the loop is replaced with a matrix command,

```

function result = diagonalSum(m)
% DIAGONALSUM returns 1 if sum of left to right diagonal of a square
% matrix is the same as that of the right to left diagonal.

[height width]=size(m);
if (width ~=height)
    error('Only works for square matrices'); % print error message and
                                              % terminate
end
leftright= sum(diag(m));
rightleft=sum(diag(m(:,width:-1:1))); % sum of diagonal of column
                                      % reversed matrix
result=(leftright==rightleft)

```

This is called “vectorization”. It is not always so easy to see how to do it. You need only worry about this if speed of your programs becomes an issue.

Exercise 11

Create a 2-dimensional data point stored in the row vector **x**. Create another 10 different 2-dimensional data points stored in the 10 by 2 matrix **A**. Create a function called `SortDist(x,A)`. The function calculates the Euclidean distance between **x** and each data point in **A**, and returns the top 3 data points in **A** that is closest to **x**. Define your function in a file in your current directory, calling it “SortDist.m”.

Note: The Euclidean distance between two 2-dimensional points $x = [x_1, x_2]$ and $y = [y_1, y_2]$ is given by $\sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2}$. Use the command “`help sort`” to explore its usage.